



C++ Style for 2014 and beyond

Bjarne Stroustrup
Morgan Stanley,
Columbia University, TAMU
www.stroustrup.com



Overview

- C++
- Simplifying code using C++14
 - For, auto, lambda, UDL, concepts, vector, move, algorithms, ...
- Overview
 - Language and libraries
- Where is the overhead?
 - Poor algorithms
 - Complicated data structures
 - Messy code



What do we want?

- Reliability
- Speed (latency)
- Throughput
- Maintainability
 - Readability
 - Productivity
 - Portability
- Popularization
 - Teaching
- C++ is my main tool
 - Research
 - Development




Language Myths

- There is a best language
 - For everybody and for every task
 - One size fits all
- Oh, no!
 - These myths confound
 - Education
 - Practice
 - Research
 - Language design
 - Management
 - Funding



What does C++ offer?

- What is C++?
 - Direct map to hardware
 - of instructions and fundamental data types
 - Initially from C
 - Zero-overhead abstraction
 - Classes with constructors and destructors, inheritance, generic programming, function objects
 - Initially from Simula
- Much of the inspiration came from operating systems
- What does C++ want to be when it grows up?
 - See above
 - And be better at it for more modern hardware and techniques



Map to Hardware

- Primitive operations => instructions

- +, %, ->, [], (), ...

value

- **int**, double, complex<double>, Date, ...

handle



- **vector**, string, thread, Matrix, ...

value

- Objects can be composed by simple concatenation:

- Arrays

- Classes/structs



- All maps to “raw memory”



Classes: Construction/Destruction

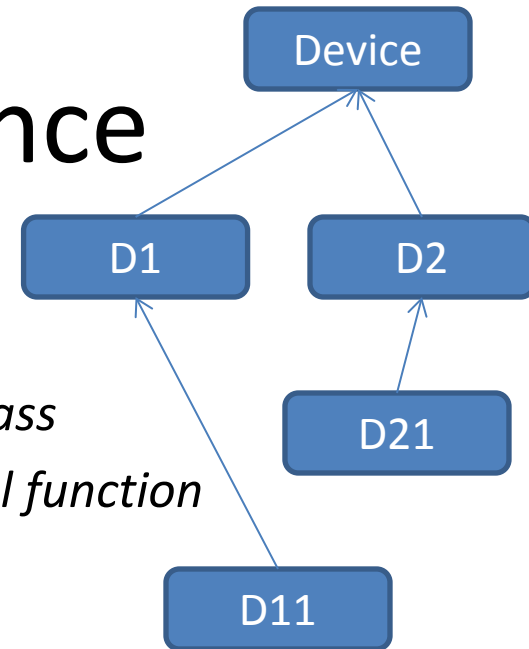
```
class X {                                // a user-defined type called X
public:                                  // interface
    X(Something);                       // constructor from Something
    ~X();                               // destructor
    // ...
private:                               // implementation
    // ...
};
```



“A constructor establishes the environment for the members to run in; the destructor reverses its actions.”

– BS 1979 (slightly rephrased)

Classes and Inheritance



- Insulate the user from the implementation

```
struct Device {                                // abstract class
    virtual int put(const char*) = 0;          // pure virtual function
    virtual int get(const char*) = 0;
};
```
- A class can be the root of a hierarchy of derived classes
 - The derived classes supply the implementation (code and possibly data)
- Abstract classes: No data members, all data in derived classes
 - “not brittle”
- Manipulate through pointer or reference
 - Typically allocated on the free store (“dynamic memory”)
 - Typically requires some form of lifetime management
 - use resource handles

Parameterized Types and Classes

- Templates
 - Essential: Support for generic programming
 - Secondary: Support for metaprogramming

template<typename T>

class vector { /* ... */ }; *// a generic type*

vector<double> constants = {3.14159265359, 2.54, 1, 6.62606957E-34, }; *// a use*

C++14 Concept

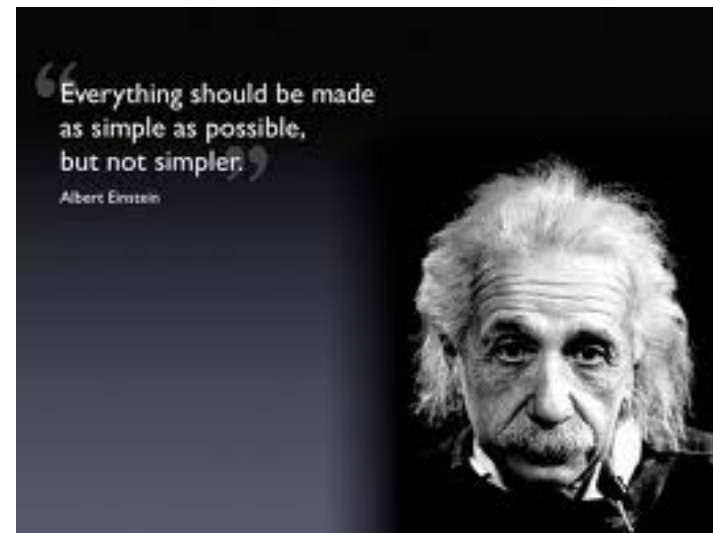
template<Sortable Seq>

void sort (Seq& c) { /* ... */ } *// a generic function taking a sortable sequence*

sort(constants); *// a use*

Make simple things simple!

- What does “simple” mean?
- Make the code directly express intent
 - Simple code is easier to understand
 - Simple code is often beautiful
 - Simple code is often fast
 - Not all code can be simple
 - Hide the complexity behind a simple interface



A simple example:

Compute the mean of a sequence of numbers

```
auto mean(const Sequence& seq)
{
    auto n = 0.0;
    for (x : seq)
        n += x;
    return n / seq.size();
}

cout << mean({1,2,3,4,5,6,7,8,9,0});

cout << char(mean("this is also a sequence"));
```

The image shows C++ code for computing the mean of a sequence. Red arrows point to the following elements: the `auto` keyword in the function signature, the `const Sequence&` parameter, the `seq` parameter, the opening curly brace of the function body, the `auto n = 0.0;` line, the `for (x : seq)` loop header, the `n += x;` loop body, the `return` statement, the closing curly brace of the function, the `mean` function call in the first `cout` statement, the sequence `{1,2,3,4,5,6,7,8,9,0}`, and the `mean` function call in the second `cout` statement.

Similar to other languages

Python

```
def mean(seq):  
    n = 0.0  
    for x in seq:  
        n += x  
    return n / len(seq)
```

C++

```
auto mean(const Sequence& seq) {  
    auto n = 0.0;  
    for (x : seq)  
        n += x;  
    return n / seq.size();  
}
```

We can simplify further

- `def mean(seq):`
 `return sum(seq) / len(seq)`
- `auto mean(const Sequence& seq) {`
 `return accumulate(seq, {}) / seq.size();`
 `}`
- Nil of the appropriate type (`Value_type<Sequence>`)
- Libraries can make most things simple to use

Resource management

- A resource is something that must be acquired and released
 - explicitly or implicitly
- Examples: memory, locks, file handles, sockets, thread handles

```
void f(int n, string name)
```

```
{
```

```
    vector<int> v(n);           // vector of n integers
```

```
    fstream fs { name, "r"};  // open file <name> for reading
```

```
    // ...
```

```
    // memory and file released here
```

- We must avoid manual resource management
 - We don't want leaks
 - We want to minimize resource retention

Value types

- The key types of C++ are value types
 - The built-in types
 - The standard-library types
 - Many of your own types
- Assignment yields two independent objects that compare equal


```
int i1 {7};  
int i2 {i1};           // i1==i2  
++i1;                  // i1==8; i2==7  
vector<int> v1 {1,2,3,4};  
vector<int> v2 {v1};    // v1==v2  
++v1[0];               // v1== {2,2,3,4}; v2=={1,2,3,4}
```

- Value types are a key to simplicity and performance

Vector representation


- Not seen by end-users (“private”)
 - **vector_rep** takes care of untyped memory
 - **vector** takes care of typed objects
 - Direct access to “raw” memory
 - C++ is (among other things) a systems programming language
 - A constructor converts raw memory to an object
 - A destructor converts an object back into raw memory

vector_rep:




Control

- We control object lifetime/life-cycle
 - Creation of objects: constructors
 - Destruction of objects: destructors
 - Movement of objects
 - Construction and assignment
 - from one scope to another
 - Copying of objects
 - Construction and assignment
 - from one scope to another
 - Access to representation
- At no cost compared to low-level hand coding




Copy and move

```
vector user(vector<int>& x)
{
    vector<int> y = x; // copy
    // ...
    return y;          // move, don't copy
}
```



```
Vector<int> z = algo(some_vec);

// copy constructor and copy assignment
// move constructor and move assignment
```




C++11/14 overview

- Language features
 - Move, uniform initialization, range-for, auto, lambdas, variadic templates, user-defined literals, digit separators, forwarding constructors, in-class initialization, generalized constant expressions, template aliases, ...
 - Memory model
- Library components
 - Type-safe threads, mutexes
 - lock-free programming
 - futures
 - Random numbers
 - Regular expressions
 - Emplace operations, move semantics, initializer lists, ...
 - Type traits

Where does the time go?

- I mean, in addition to
 - System calls
 - Net accesses
 - Disk read/writesof course.
- A partial answer
 - Too much data
 - Poorly structured data
 - Unpredictable access to data
 - Messy code
 - Too many run-time decisions



Use compact data


- Vector vs. list
- Object layout



Vector vs. List

- Generate N random integers and insert them into a sequence so that each is inserted in its proper position in the numerical order. **5 1 4 2** gives:
 - **5**
 - **1 5**
 - **1 4 5**
 - **1 2 4 5**
- Remove elements one at a time by picking a random position in the sequence and removing the element there. Positions **1 2 0 0** gives
 - **1 2 4 5**
 - **1 4 5**
 - **1 4**
 - **4**
- For which N is it better to use a linked list than a vector (or an array) to represent the sequence?
- The sequence grows incrementally

Vector vs. List



- Vector beats list massively for insertion and deletion
 - For small elements and relatively small numbers (up to 500,000 on my machine)
 - Your mileage **will** vary

Implementation (performance-critical part)


```
template<class C> void do_insert(C& s, int N)
{
    for (int i=0; i<N; ++i) insert(s,randval[i]);
}
```

```
template<class C> void do_erase(C& s, int N)
{
    for (int i=0; i<N; ++i) {
        auto p = s.begin();
        int count = randval2[i];
        while (count-->0) ++p;    //advance(p,randval2[i]) would optimize vector version
        s.erase(p);              // remove element at position N
    }
}
```

```
template<class C> void insert(C& s, int n)
{
    auto p = find_if(s.begin(),s.end(),[n](int i){ return i>n;}); // find first larger or end
    s.insert(p,n);
}
```

Vector vs. List

- The amount of memory used differ dramatically
 - List uses 4+ words per element
 - Vector uses 1 word per element
- An unoptimized list does more allocations
 - One per element
 - Can be eliminated by pre-allocation
- Memory access is relatively slow
 - Unpredictable memory access gives many more cache misses
- Find the insertion and deletion points
 - Linear search
- Use a map?
 - But then we could use binary search and indexed access for a vector
 - Also: this misses the point. Don't just optimize for the exercise




Vector vs. List

- Implications:

- Don't store data unnecessarily.
- Keep data compact.
- Access memory in a predictable manner.

- Measure!

- Your intuition is not reliable
- Complexity theory is only a **very** rough guide



Experiment!


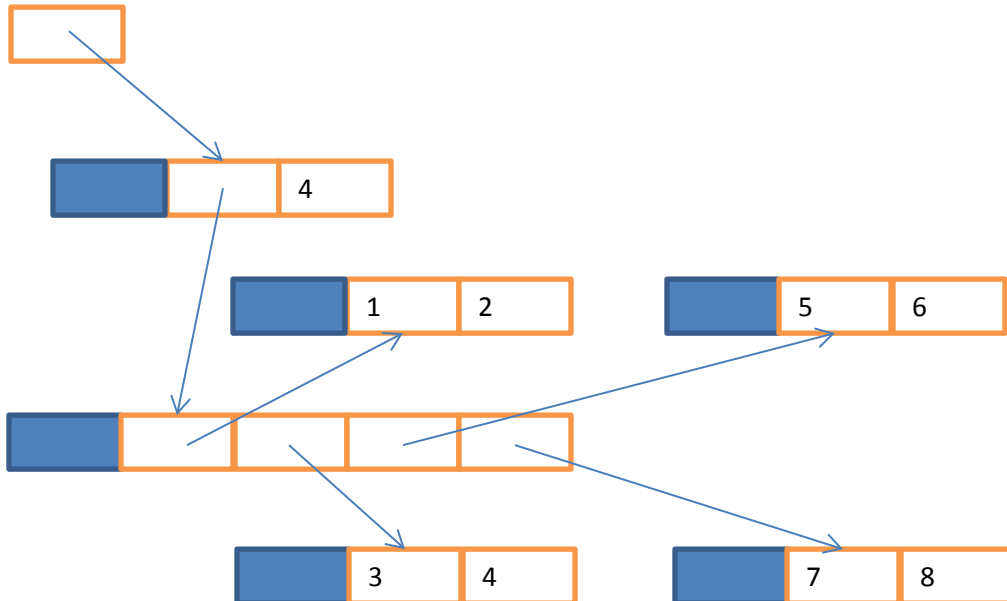
- Measure
- Control
- Simplify
- Reason



Use compact layout

- `vector<Point> vp = { Point{1,2}, Point{3,4}, Point{5,6}, Point{7,8} };`

“True OO” style:



C++98 to C++14

C++98:

```
circle* p = new circle(pnt,42);  
vector<shape*> v = load_shapes();  
for (vector<shape*>::iterator i = v.begin(); i != v.end(); ++i) {  
    if (*i && **i == *p)  
        cout << **i << " is a match\n";  
}
```

not exception-safe
missing try/catch,

// ... later, possibly elsewhere ...

```
for( vector<shape*>::iterator i = v.begin(); i != v.end(); ++i  
) {  
    delete *i;  
}  
delete p;
```

Easy to forget;
easy to get wrong

C++14:

```
auto p = make_unique<circle>(pnt,42);  
auto v = load_shapes();  
for (s : v) {  
    if (s && *s == *p)  
        cout << *s << " is a match\n";  
}
```


C++11/14

- It feels like a new language
 - I find it easier to express my ideas
 - My code is clearer and more compact
 - My programs are faster
- You get few benefits if you insist
 - Writing in 1970s C style
 - Writing in 1980s “Pure OO” style



Uniform initialization

- You can use `{}`-initialization for all types in all contexts

```
int a[] = { 1,2,3 };
```

```
vector<int> v { 1,2,3 };
```

```
vector<string> geek_heros = {  
    "Dahl", "Kernighan", "McIlroy", "Nygaard ", "Ritchie", "Stepanov"  
};
```

```
thread t {}; // default initialization
```

// remember “thread t();” is a function declaration

```
complex<double> z {1,2}; // invokes constructor
```

```
struct S { double x, y; } s {1,2}; // no constructor (just initialize members)
```

Uniform initialization

- `{}`-initialization `X{v}` yields the same value of `x` in every context

```
X x{a};
```

```
X* p = new X{a};
```

```
z = X{a};           // use as cast
```

```
void f(X);
```

```
f({a});             // function argument (of type X)
```

```
X g() {
```

```
    // ...
```

```
    return {a};      // function return value (function returning X)
```

```
}
```

```
Y::Y(a) : X{a} { /* ... */ };    // base class initializer
```

auto

- Deduce a type of an object from its initializer

```
auto x = 1;           // x is an int
auto y = 1.2;         // y is a double
```

- Most useful when types gets hard to type or hard to know

```
template<class C>
```

```
void use(C& c)
```

```
{
```

```
    for (auto p = c.begin(); p!=c.end(); ++p)    // p is a ???
```

```
        cout << *p << '\n';
```

```
}
```

- Curio: The oldest C++11 feature
 - I implemented it in 1983/84

range-for

- Make the simplest loops simpler

```
template<class C>
void use(C& c)
{
    for (auto x : c)
        cout << x << '\n';
}

for(auto x : { 1, 2, 5, 8, 13})
    test(x);
```



User-Defined Literals

- Examples
 - `"Hello! "` `// const char*`
 - `"Howdy! "s` `// std::string`
 - `2.3*5.7i` `// "i" for "imaginary": a complex number`
 - `4h+6min+3s` `// 4 hours, 6 minutes, and 3 seconds`
- Can be used for type-rich programming
 - `Speed s = 100m/9s;` `// very fast for a human`
 - `Acceleration a1 = s/9s;` `// OK`
 - `Acceleration a2 = s;` `// error: unit mismatch`
- Definition
 - `complex<double> operator "" i(long double d) { return {0,d}; }`

General constant expressions

- Think
 - ROM
 - concurrency
 - Compile-time computation (performance, compactness)
 - Type safety (reliability, maintainability)

```
constexpr int abs(int i) { return (0<=i) ? i : -i; } // can be constant expression
```

```
struct Point {  
    int x, y;  
    constexpr Point(int xx, int yy) : x{xx}, y{yy} { } // "literal type"  
};
```

```
constexpr Point p1{1,2}; // must be evaluated at compile time: ok  
constexpr Point p2{p1.y,abs(x)}; // ok?: is x is a constant expression?
```


Lambda expressions

- A lambda expression (“a lambda”) is a use-once function object

```
template<class C, class Oper>
```

```
void for_all(C& c, Oper op)           // assume that C is a container of pointers
```

```
{
```

```
    for (auto& x : c)
```

```
        op(*x);    // pass op() a reference to each element pointed to
```

```
}
```

```
void user()
```

```
{
```

```
    vector<unique_ptr<Shape>> v;
```

```
    while (cin)
```

```
        v.push_back(read_shape(cin));    // read shape from input
```

```
    for_all(v, [](Shape& s){ s.draw(); });
```

```
// draw_all()
```

```
    for_all(v, [](Shape& s){ s.rotate(45); });
```

```
// rotate_all(45)
```

```
}
```

Variadic templates

- Any number of arguments of any types

```
template <class F, class ...Args>           // thread constructor  
    explicit thread(F&& f, Args&&... args); // argument types must  
                                              // match the operation's  
                                              // argument types
```

```
void f0();           // no arguments  
void f1(int);        // one int argument
```

```
thread t1 {f0};  
thread t2 {f0,1};           // error: too many arguments  
thread t3 {f1};             // error: too few arguments  
thread t4 {f1,1};  
thread t5 {f1,1,2};         // error: too many arguments  
thread t3 {f1,"I'm being silly"}; // error: wrong type of argument
```


Template aliases

- Notation matters
- C++98 exposes all details when we use templates
`typename iterator_traits<For>::value_type x;`
- C++11 allows us to hide details
`template<typename Iter>
using Value_type<T> = typename std::iterator_traits<For>::value_type;
// ...
Value_type<For> x;`
- Had I had an initializer, I could have used **auto**
`auto x = *p;`




Programming
Principles and Practice Using C++

C++ Information




- www.isocpp.org
 - The C++ Foundation's website
 - Standards information, articles, user-group information
- Bjarne Stroustrup
 - *A Tour of C++*: All of C++ in 180 pages
 - *The C++ Programming Language* (4th edition): All of C++ in 1,300 pages
 - *Programming: Principles and Practice using C++* (2nd edition)
 - www.stroustrup.com: Publication list, C++ libraries, FAQs, etc.
- The ISO Standards Committee site
 - Search for “WG21”
 - *The ISO standard*: All of C++ in 1,300 pages of “standardese”
 - All committee documents (incl. proposals)



Questions?

C++: A light-weight abstraction programming language



Key strengths:

- software infrastructure
- resource-constrained applications

Practice type-rich programming